

# Running Apache Spark & Apache Zeppelin in Production

**Vinay Shukla**

Director, Product Management

August 31, 2016

Twitter: @neomythos



# Who am i?

## Vinay Shukla

- Product Management
- Spark for 2.5 + years, Hadoop for 3+ years
- Recovering Programmer
- Blog at [www.vinayshukla.com](http://www.vinayshukla.com)
- Twitter: @neomythos
- Addicted to Yoga, Hiking, & Coffee
- Smallest contributor to Apache Zeppelin

# Security: Rings of Defense

## Perimeter Level Security

- Network Security (i.e. Firewalls)

## Authentication

- Kerberos
- Knox (Other Gateways)

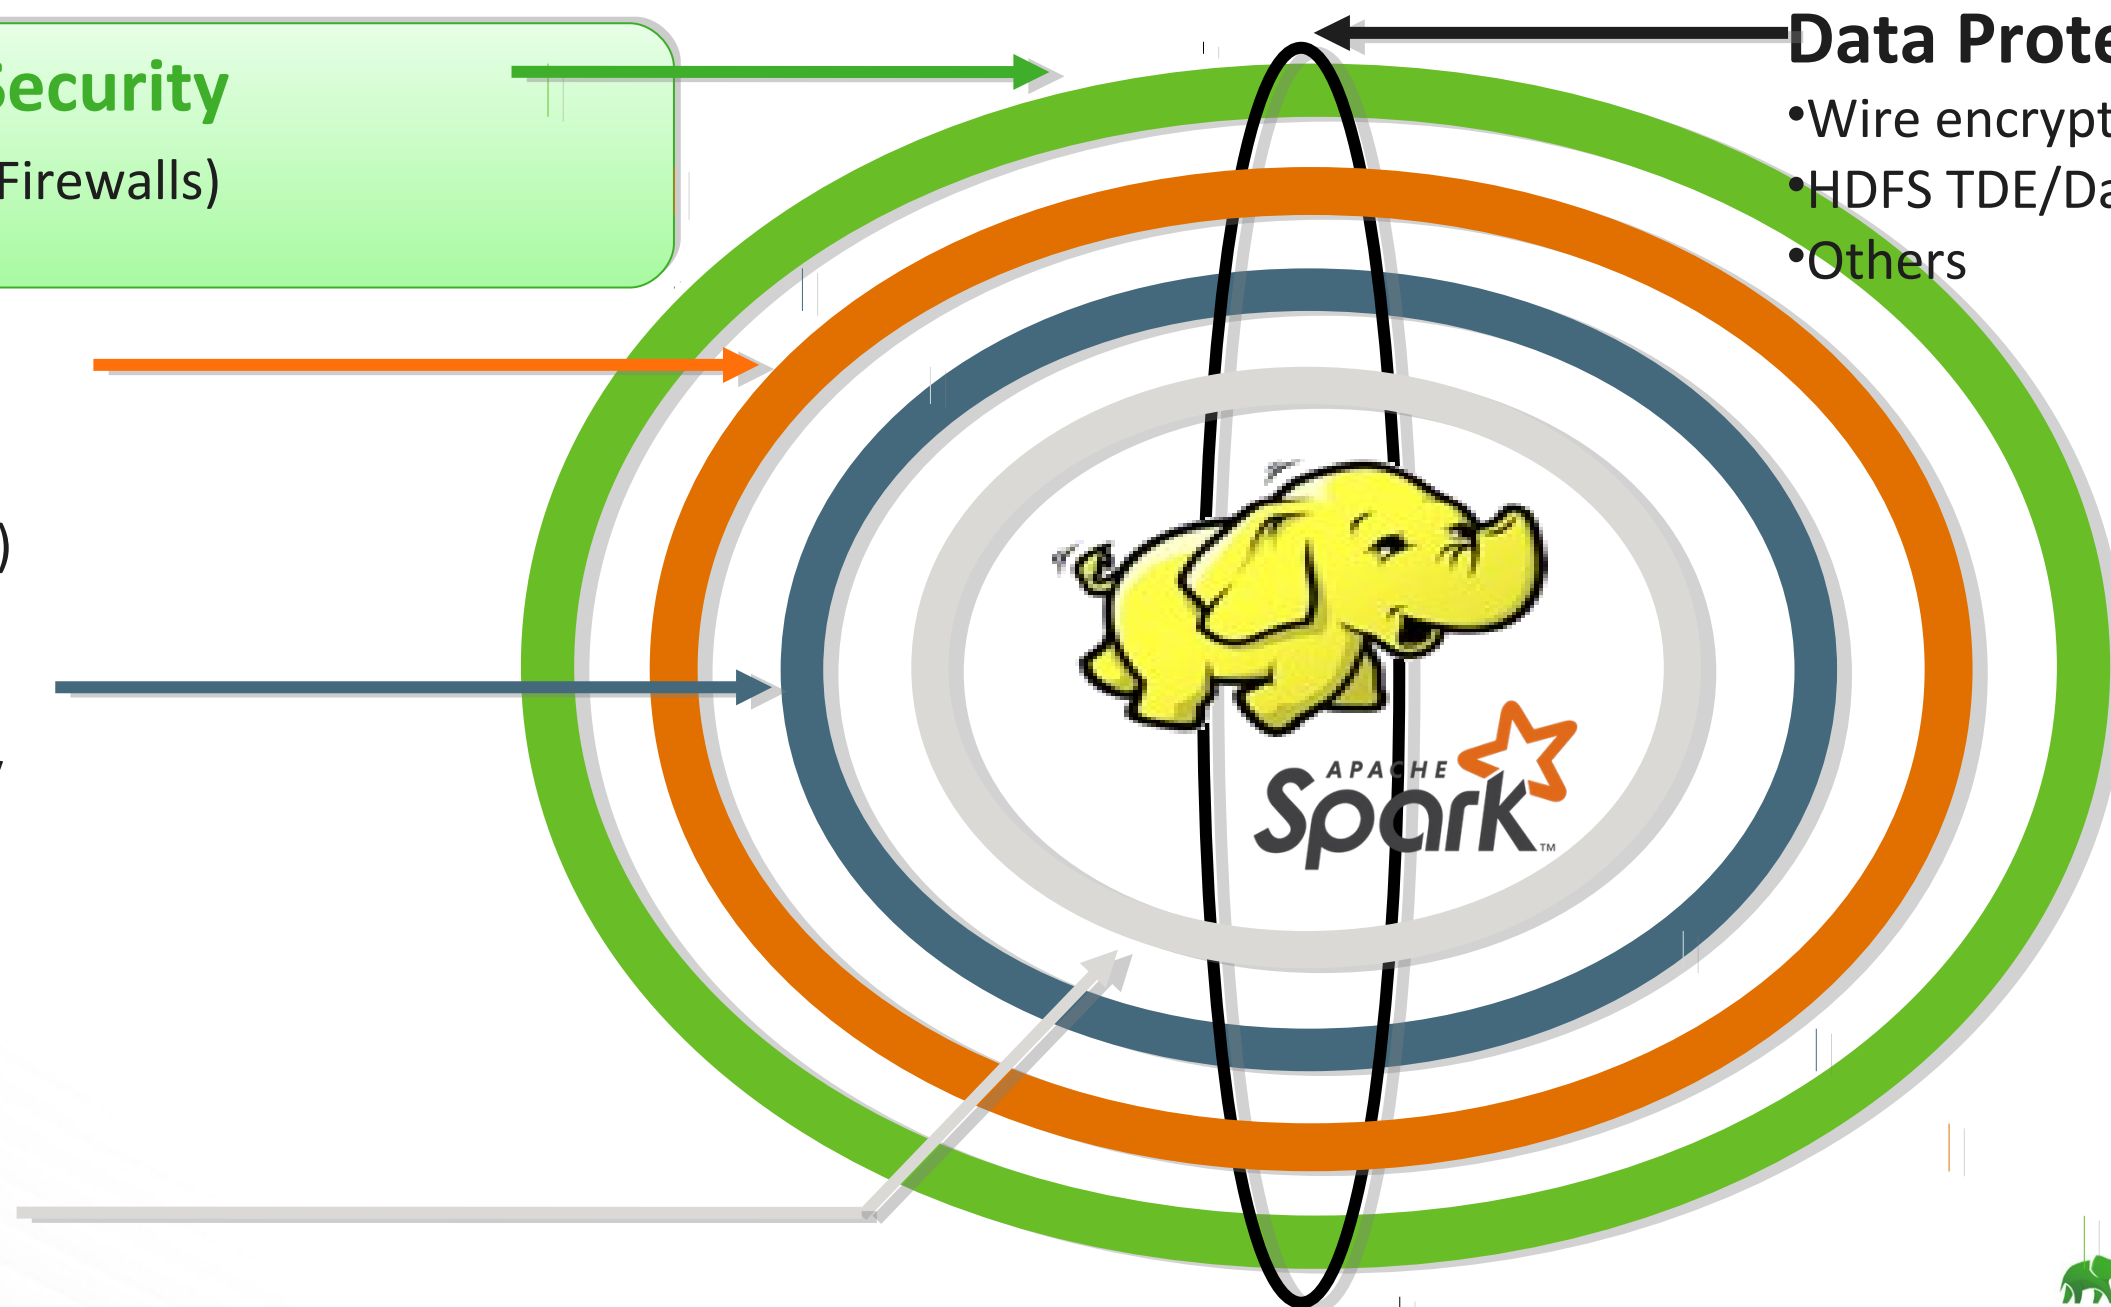
## Authorization

- Apache Ranger/Sentry
- HDFS Permissions
- HDFS ACLs
- YARN ACL

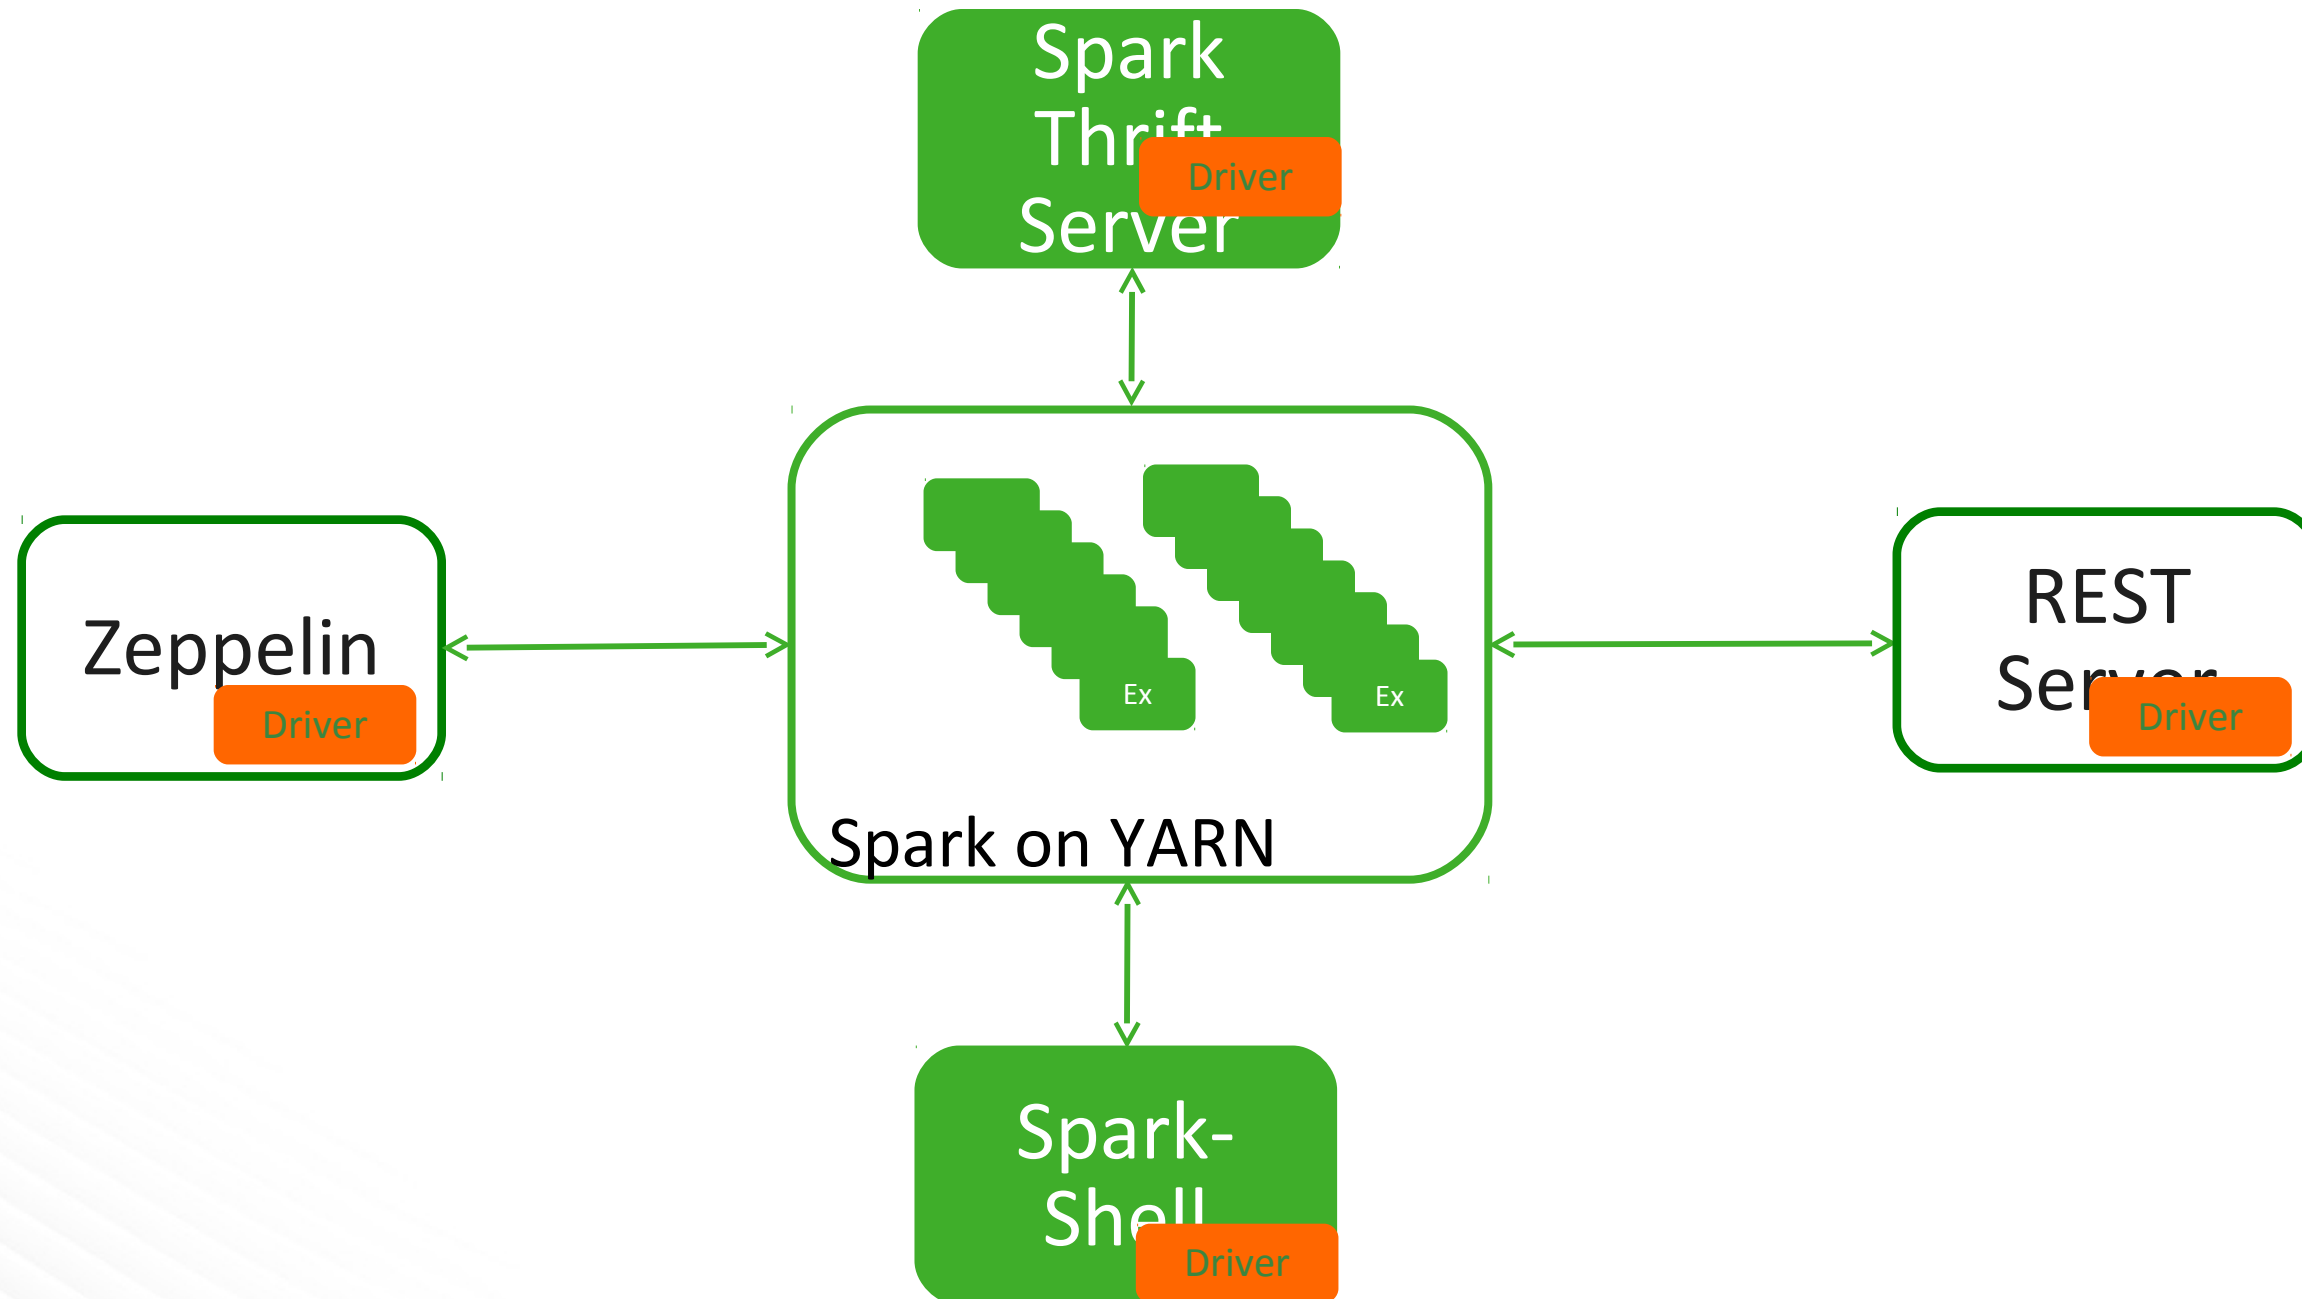
## OS Security

## Data Protection

- Wire encryption
- HDFS TDE/Dare
- Others

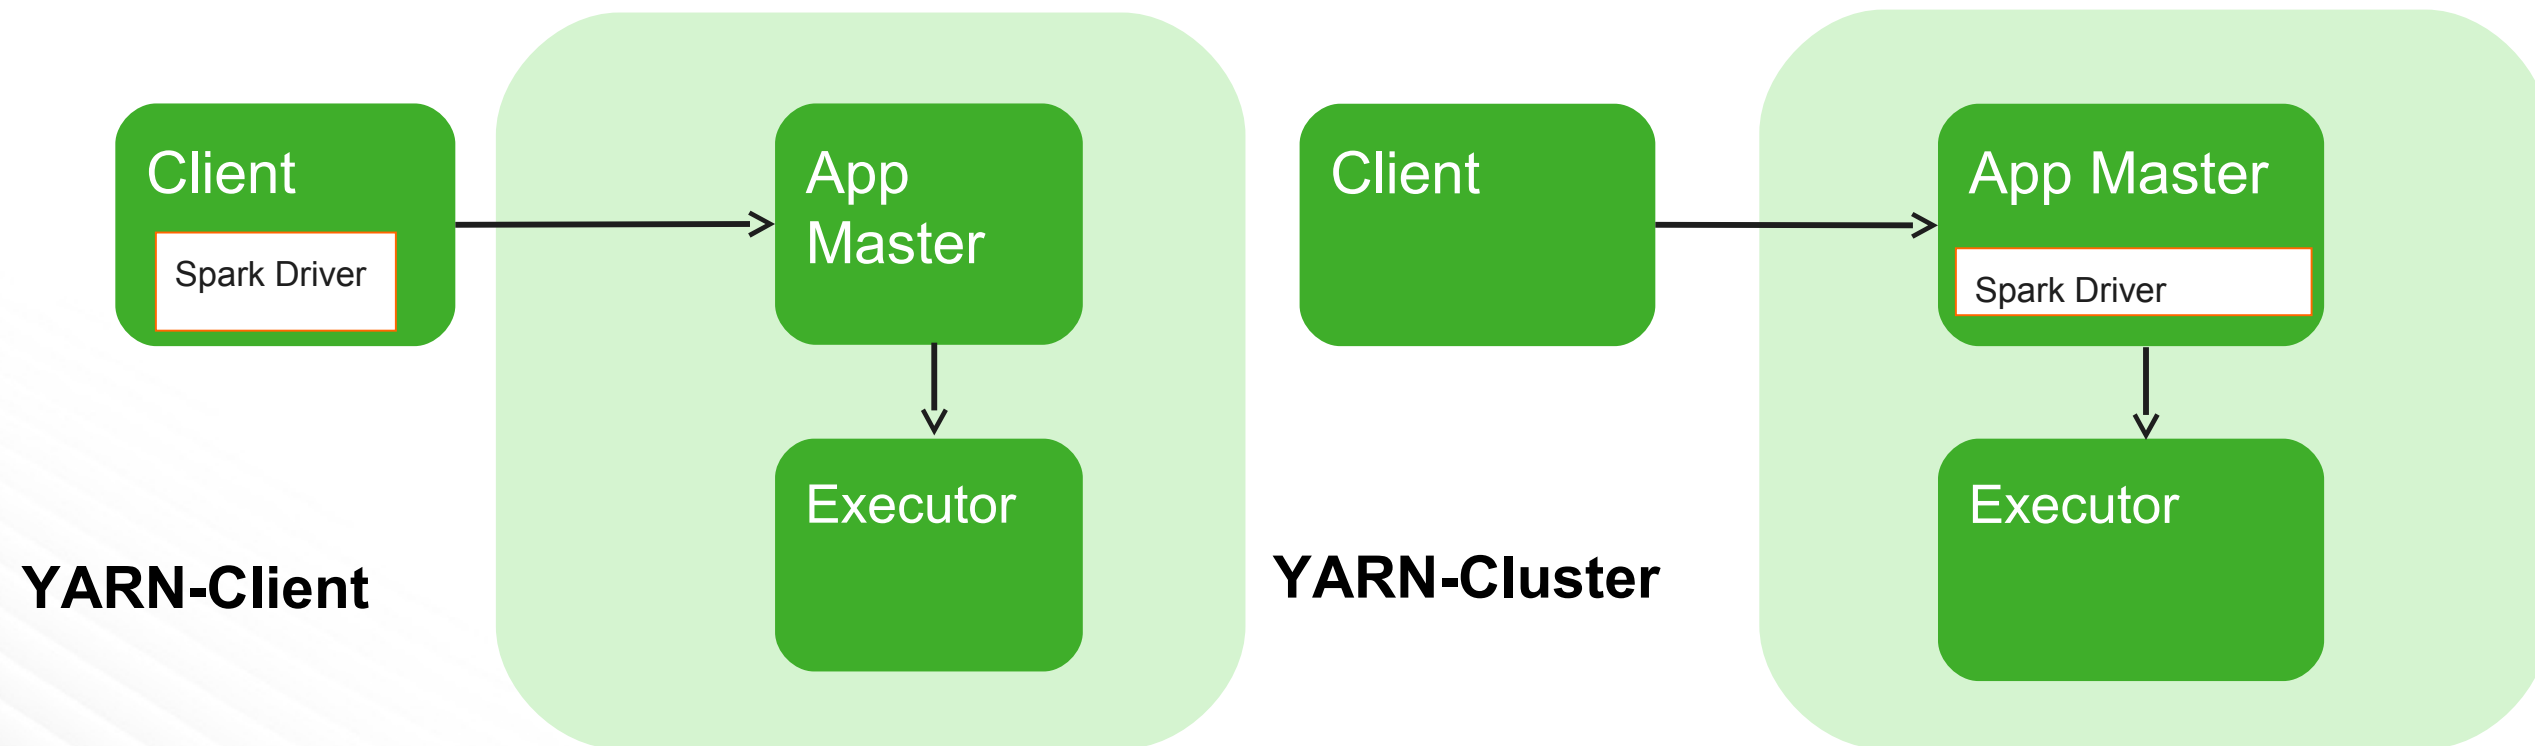


# Interacting with Spark

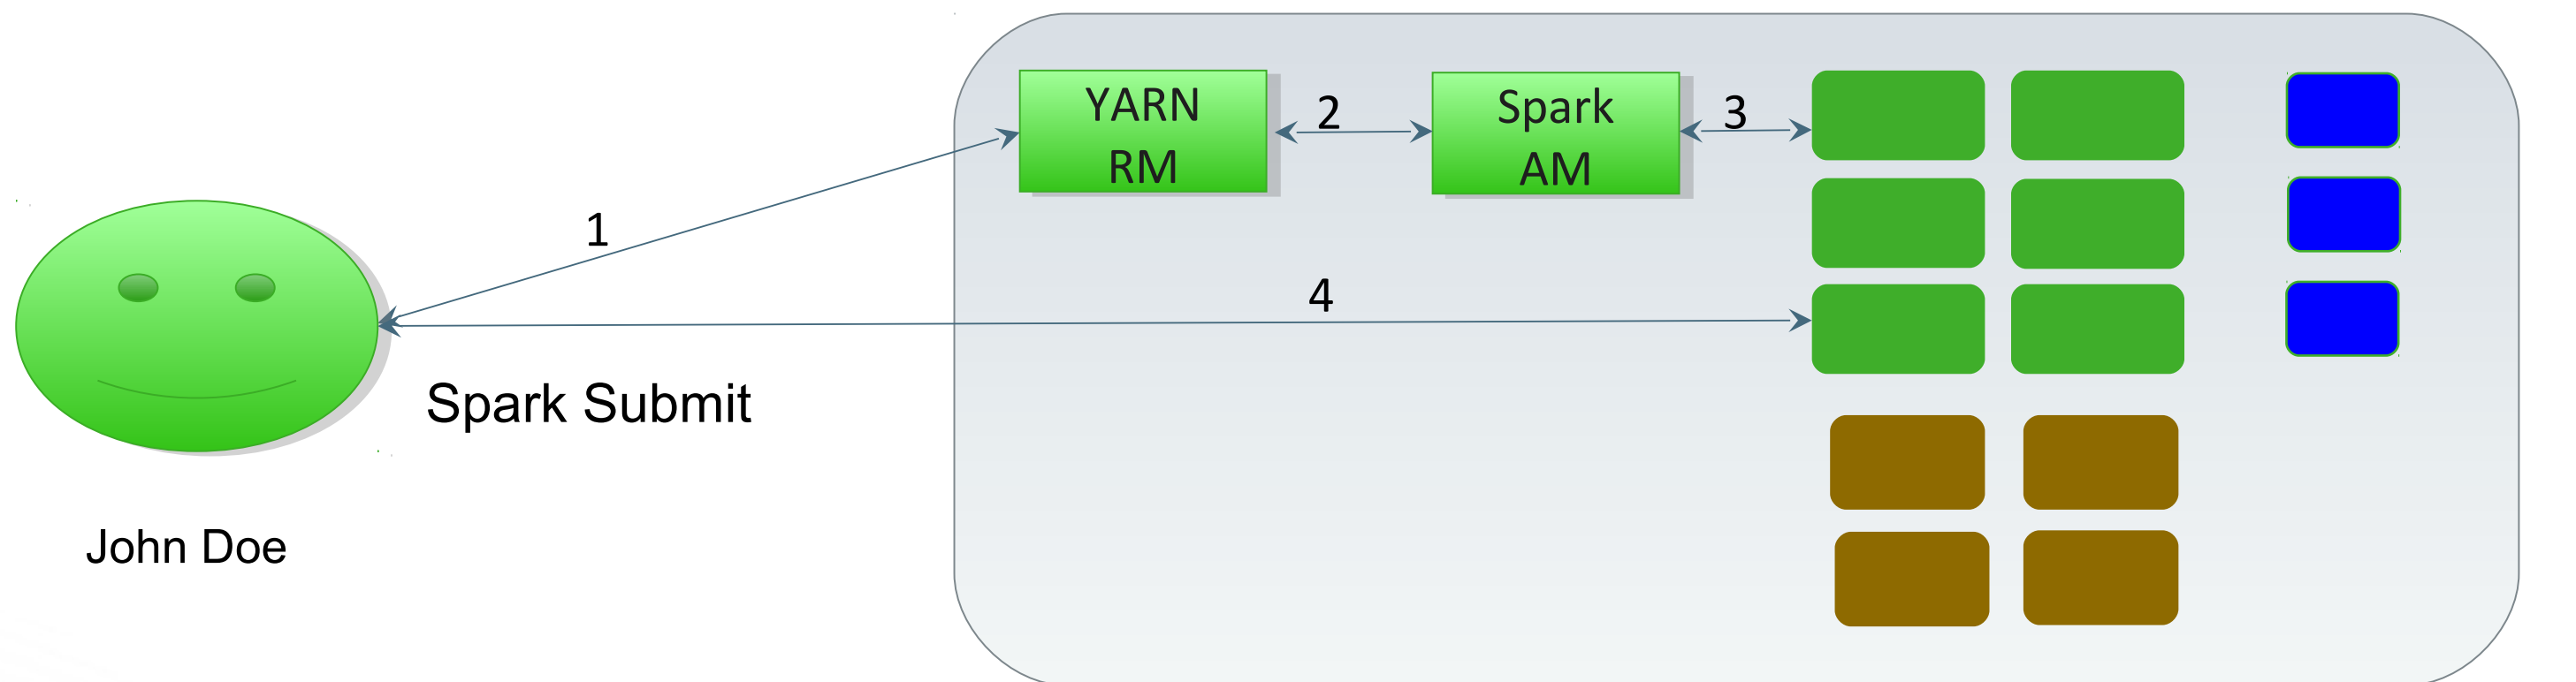


# Context: Spark Deployment Modes

- Spark on YARN
  - Spark driver (SparkContext) in YARN AM(yarn-cluster)
  - Spark driver (SparkContext) in local (yarn-client):
    - Spark Shell & Spark Thrift Server runs in yarn-client only



# Spark on YARN



John Doe

Spark Submit

**Hadoop Cluster**



HDFS



Node  
Manager



Executor





# Spark – Security – Four Pillars

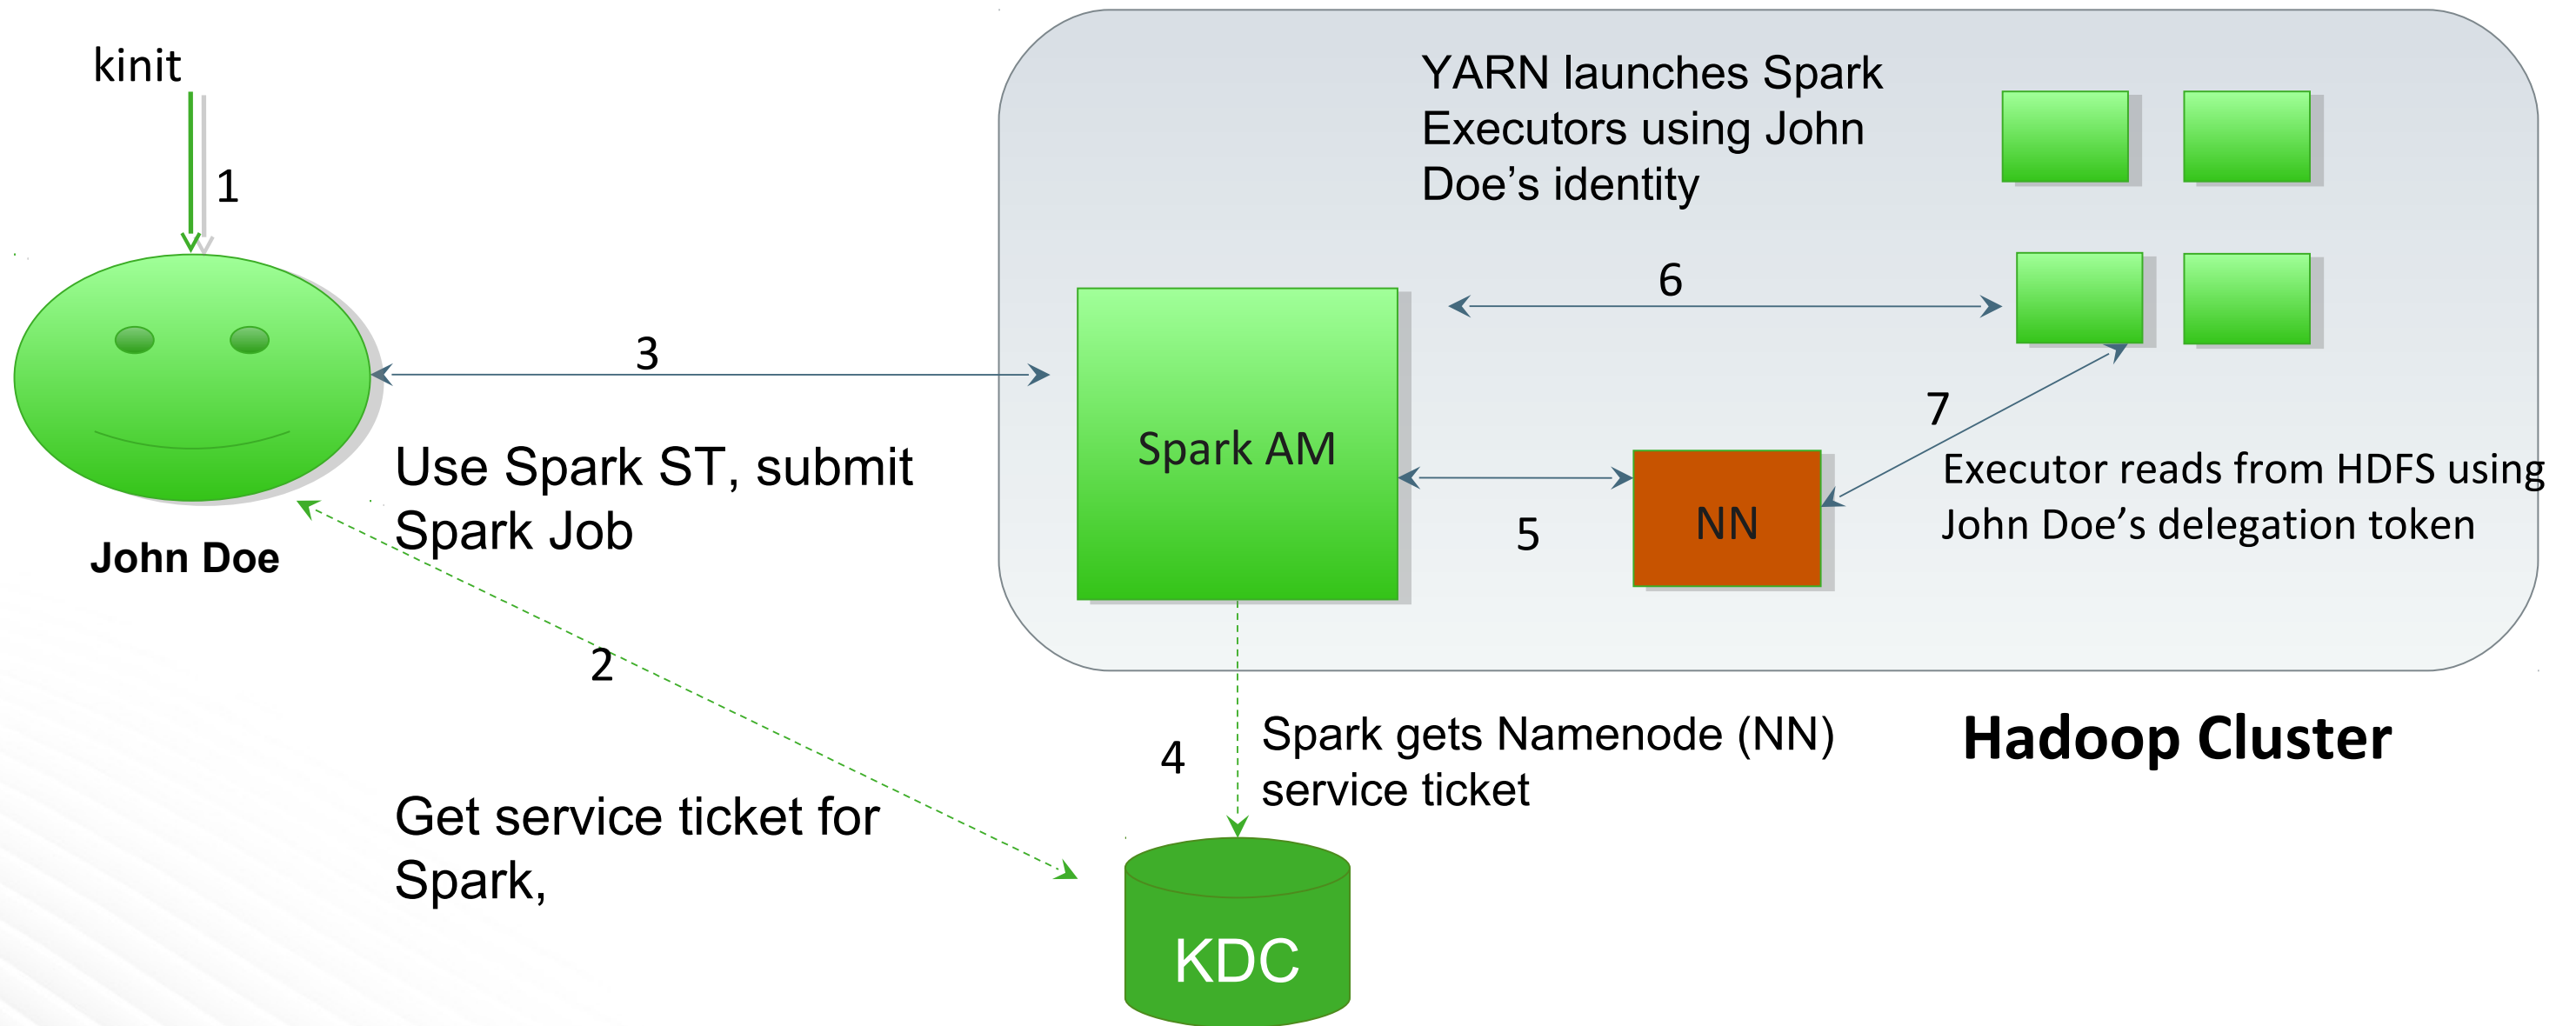
- Authentication
- Authorization
- Audit
- Encryption

Ensure network is secure



Spark leverages Kerberos on YARN

# Kerberos authentication within Spark



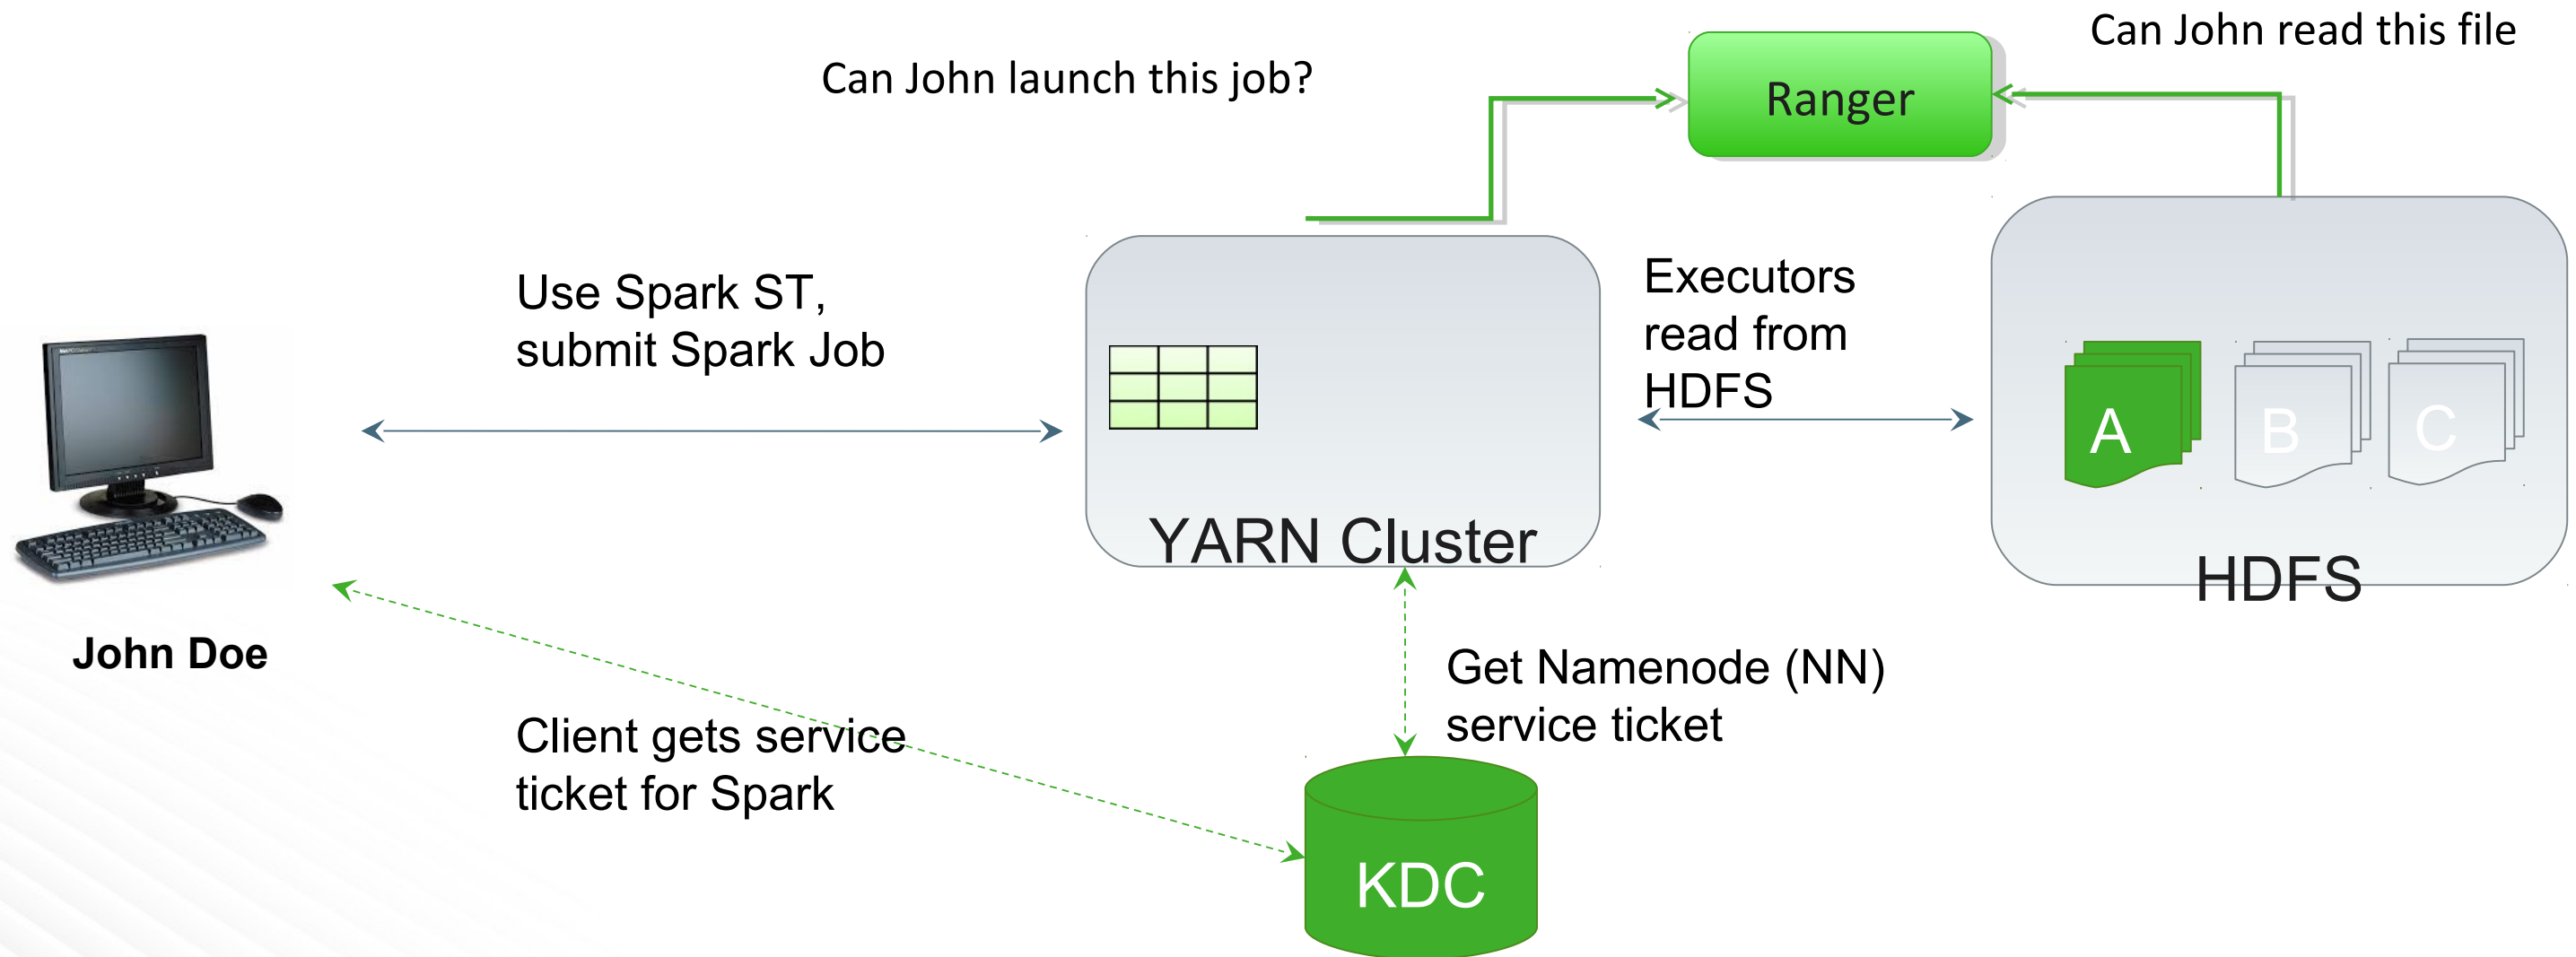


# Spark – Kerberos - Example

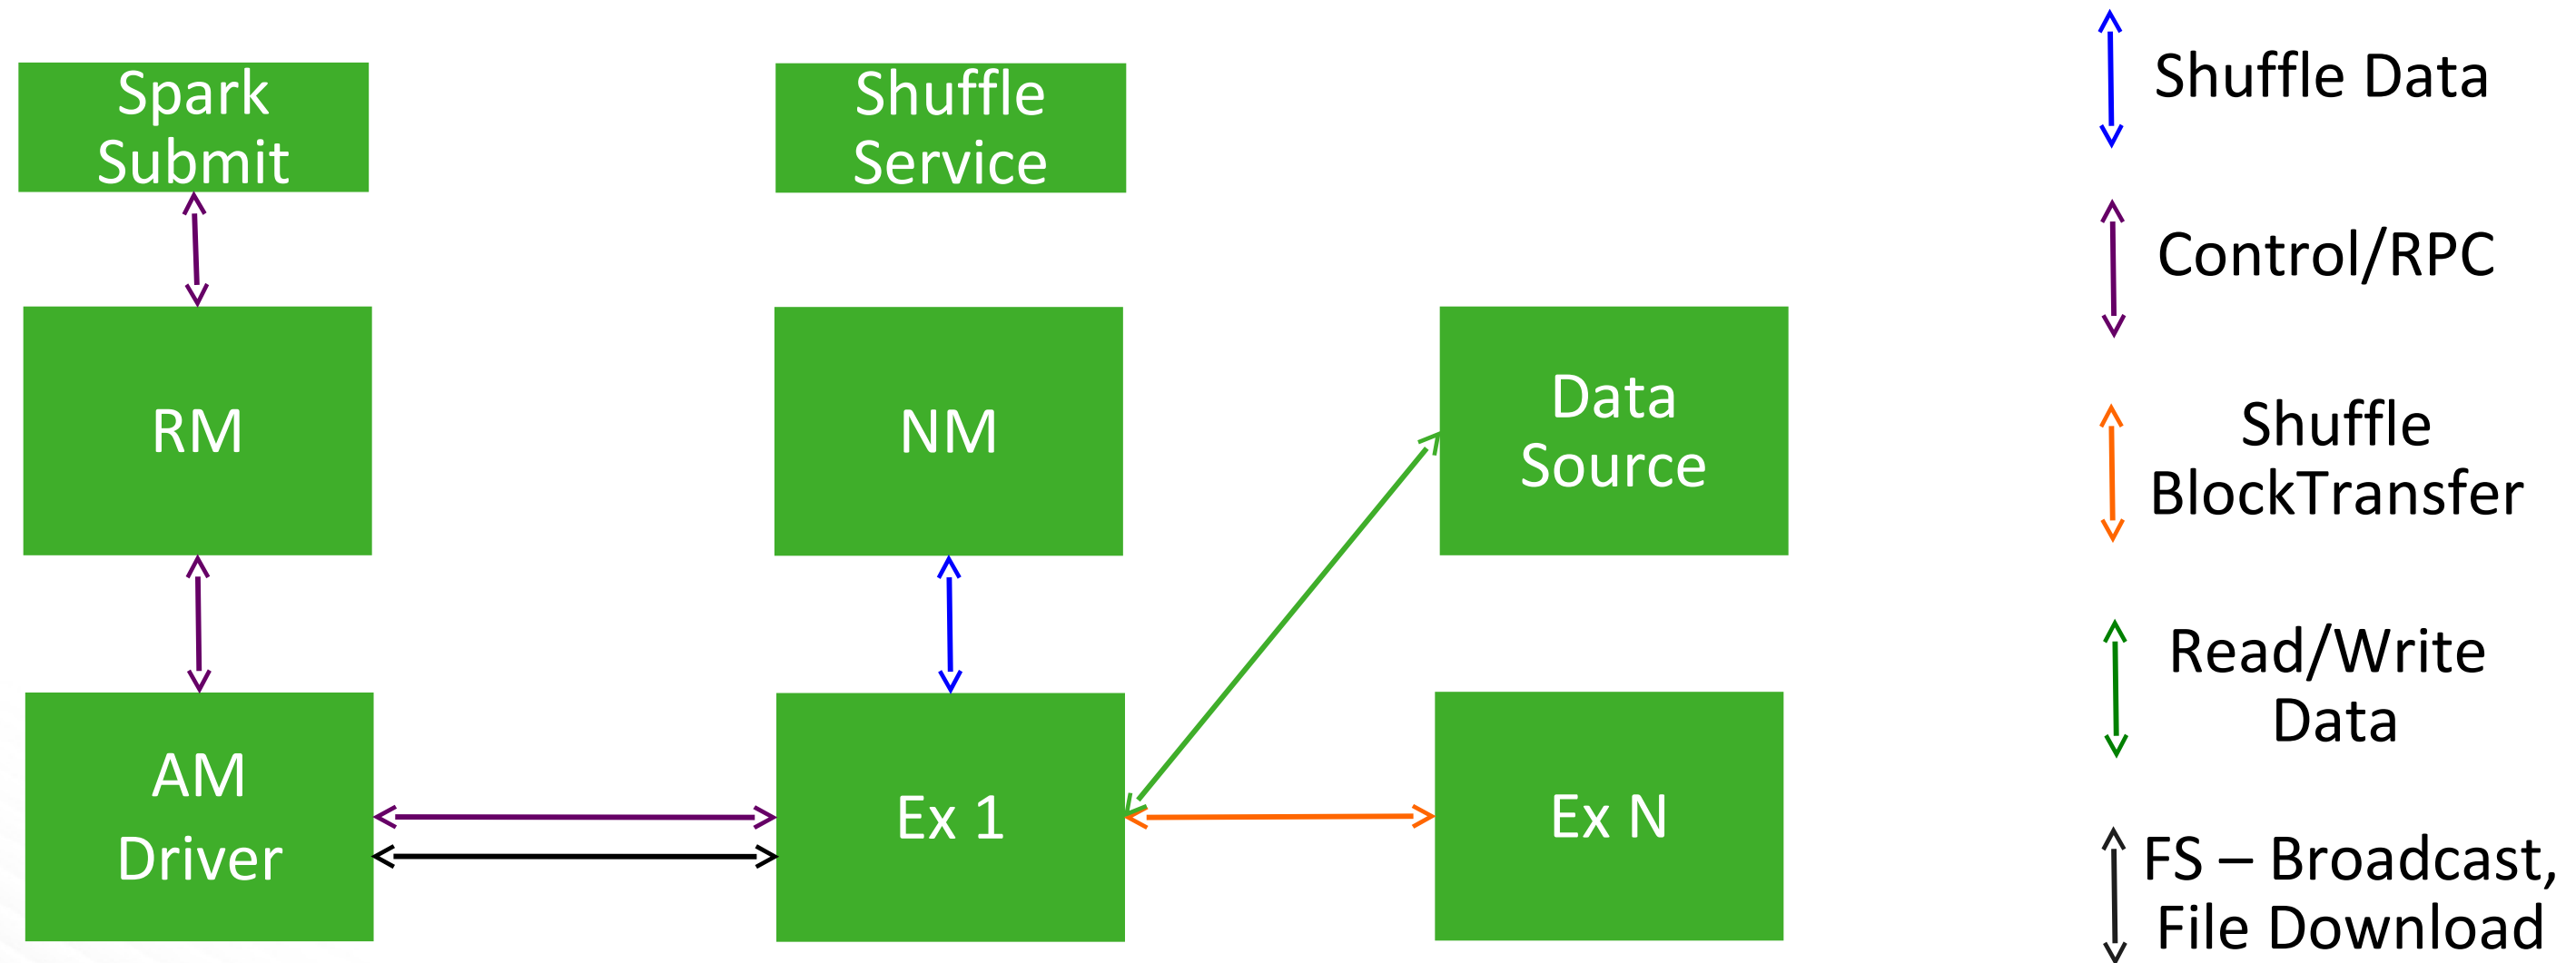
```
kinit -kt /etc/security/keytabs/johndoe.keytab johndoe@  
EXAMPLE.COM
```

```
./bin/spark-submit --class org.apache.spark.examples.SparkPi  
--master yarn-cluster --num-executors 3 --driver-memory 512m  
--executor-memory 512m --executor-cores 1 lib/spark-  
examples*.jar 10
```

# Spark – Authorization



# Encryption: Spark – Communication Channels



# Spark Communication Encryption Settings

↕	Shuffle Data	NM > Ex leverages YARN based SSL
↕	Control/RPC	spark.authenticate = true. Leverage YARN to distribute keys
↕	Shuffle BlockTransfer	spark.authenticate.enableSaslEncryption= true
↕	Read/Write Data	Depends on Data Source, For HDFS RPC (RC4   3DES) or SSL for WebHDFS
↕	FS – Broadcast, File Download	spark.ssl.enabled = true

# Sharp Edges with Spark Security

- ❑ SparkSQL – Only coarse grain access control today
- ❑ Client -> Spark Thrift Server > Spark Executors – No identity propagation on 2<sup>nd</sup> hop
  - Lowers security, forces STS to run as Hive user to read all data
  - Use SparkSQL via shell or programmatic API
  - <https://issues.apache.org/jira/browse/SPARK-5159>
- ❑ Spark Stream + Kafka + Kerberos
  - Issues fixed in HDP 2.4.x
  - No SSL support yet
- ❑ Spark Shuffle > Only SASL, no SSL support
- ❑ Spark Shuffle > No encryption for spill to disk or intermediate data

Fine grained Security to  
SparkSQL

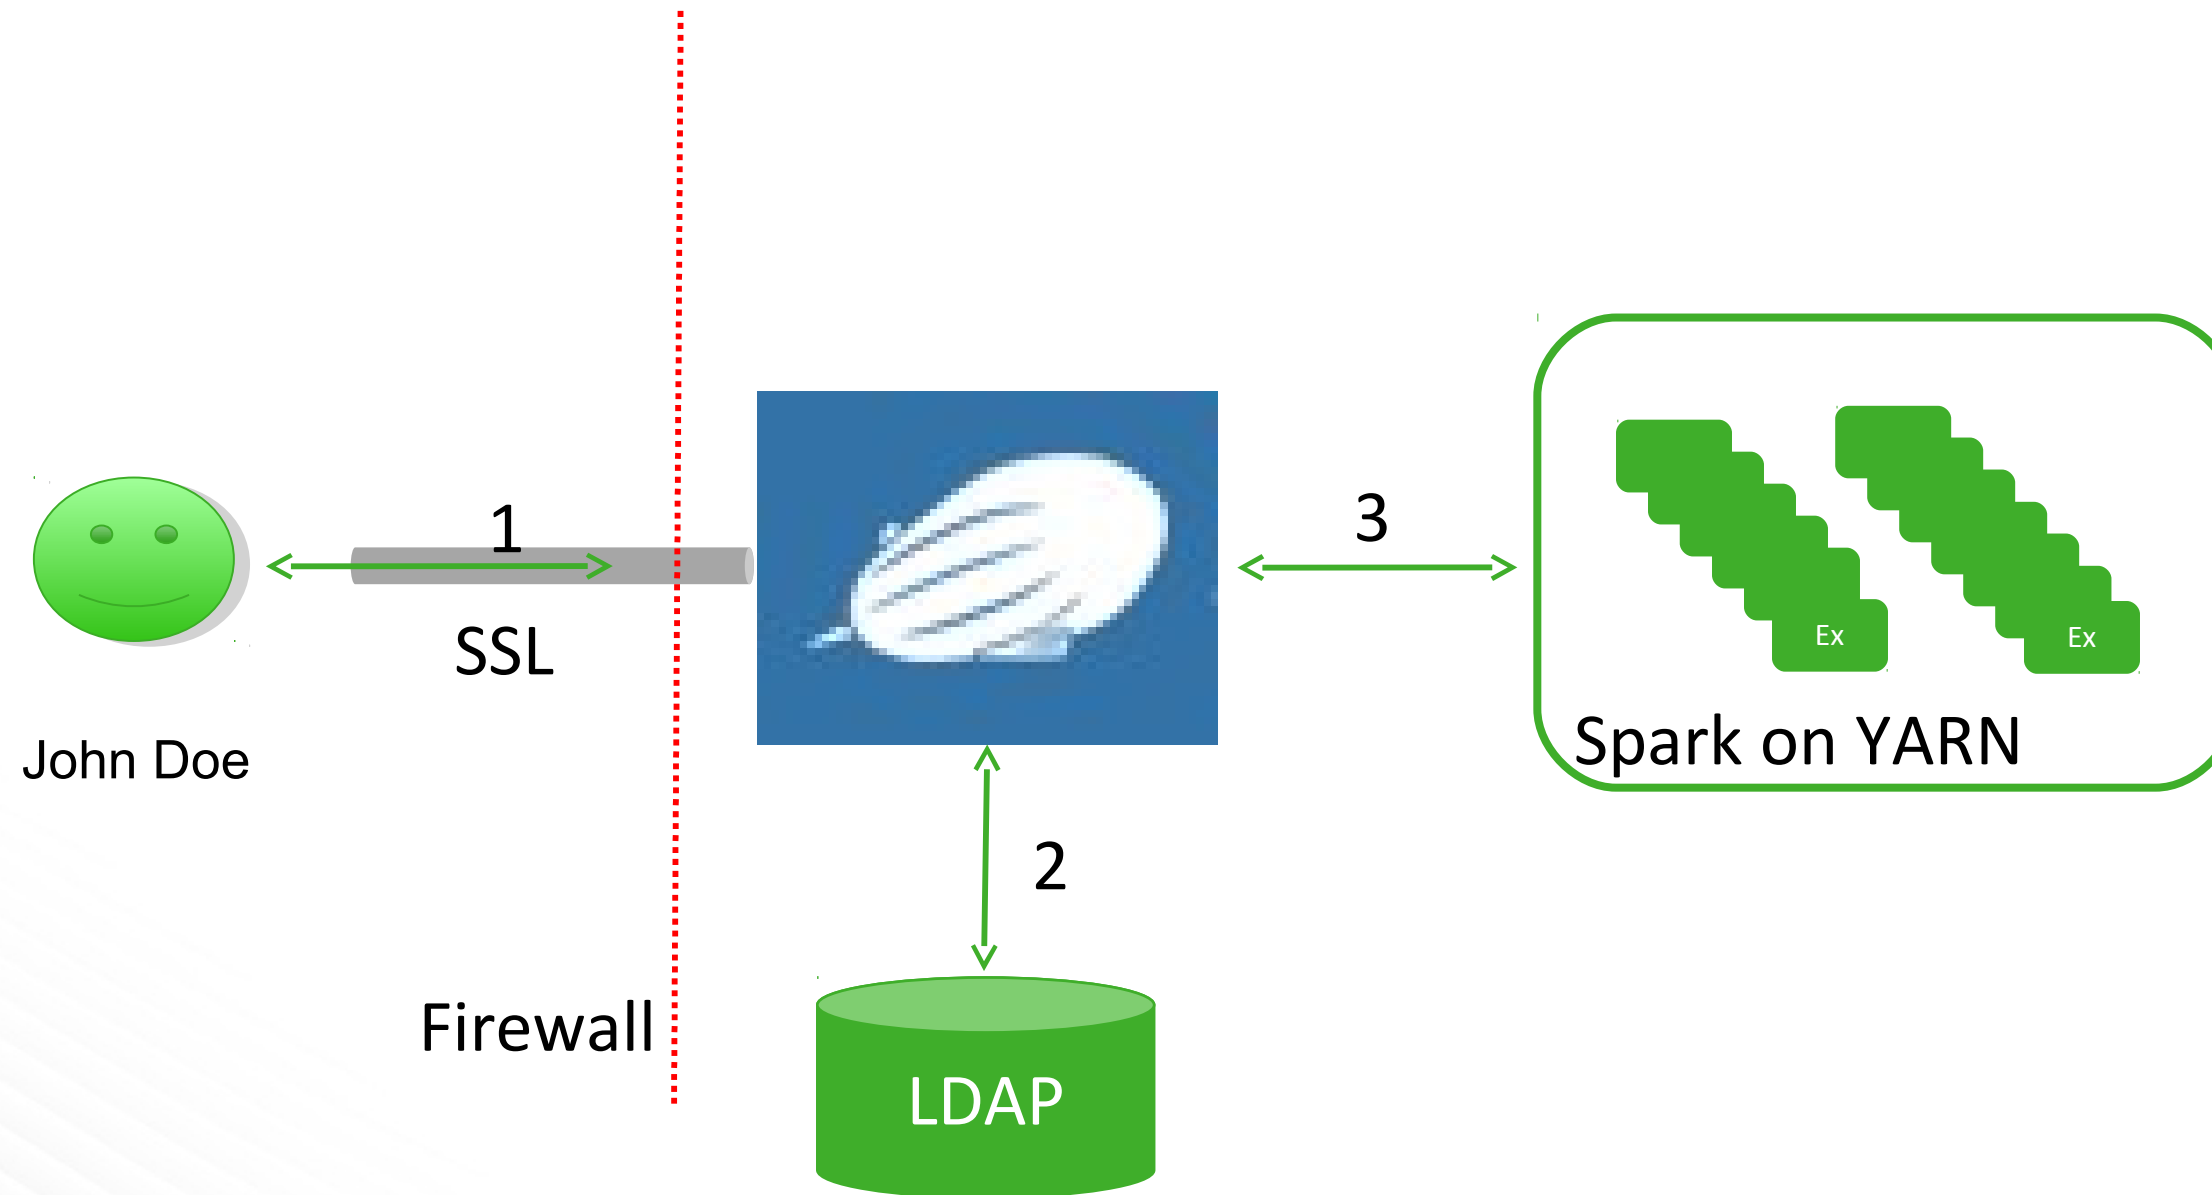
<http://bit.ly/2bLghGz>

<http://bit.ly/2bTX7Pm>

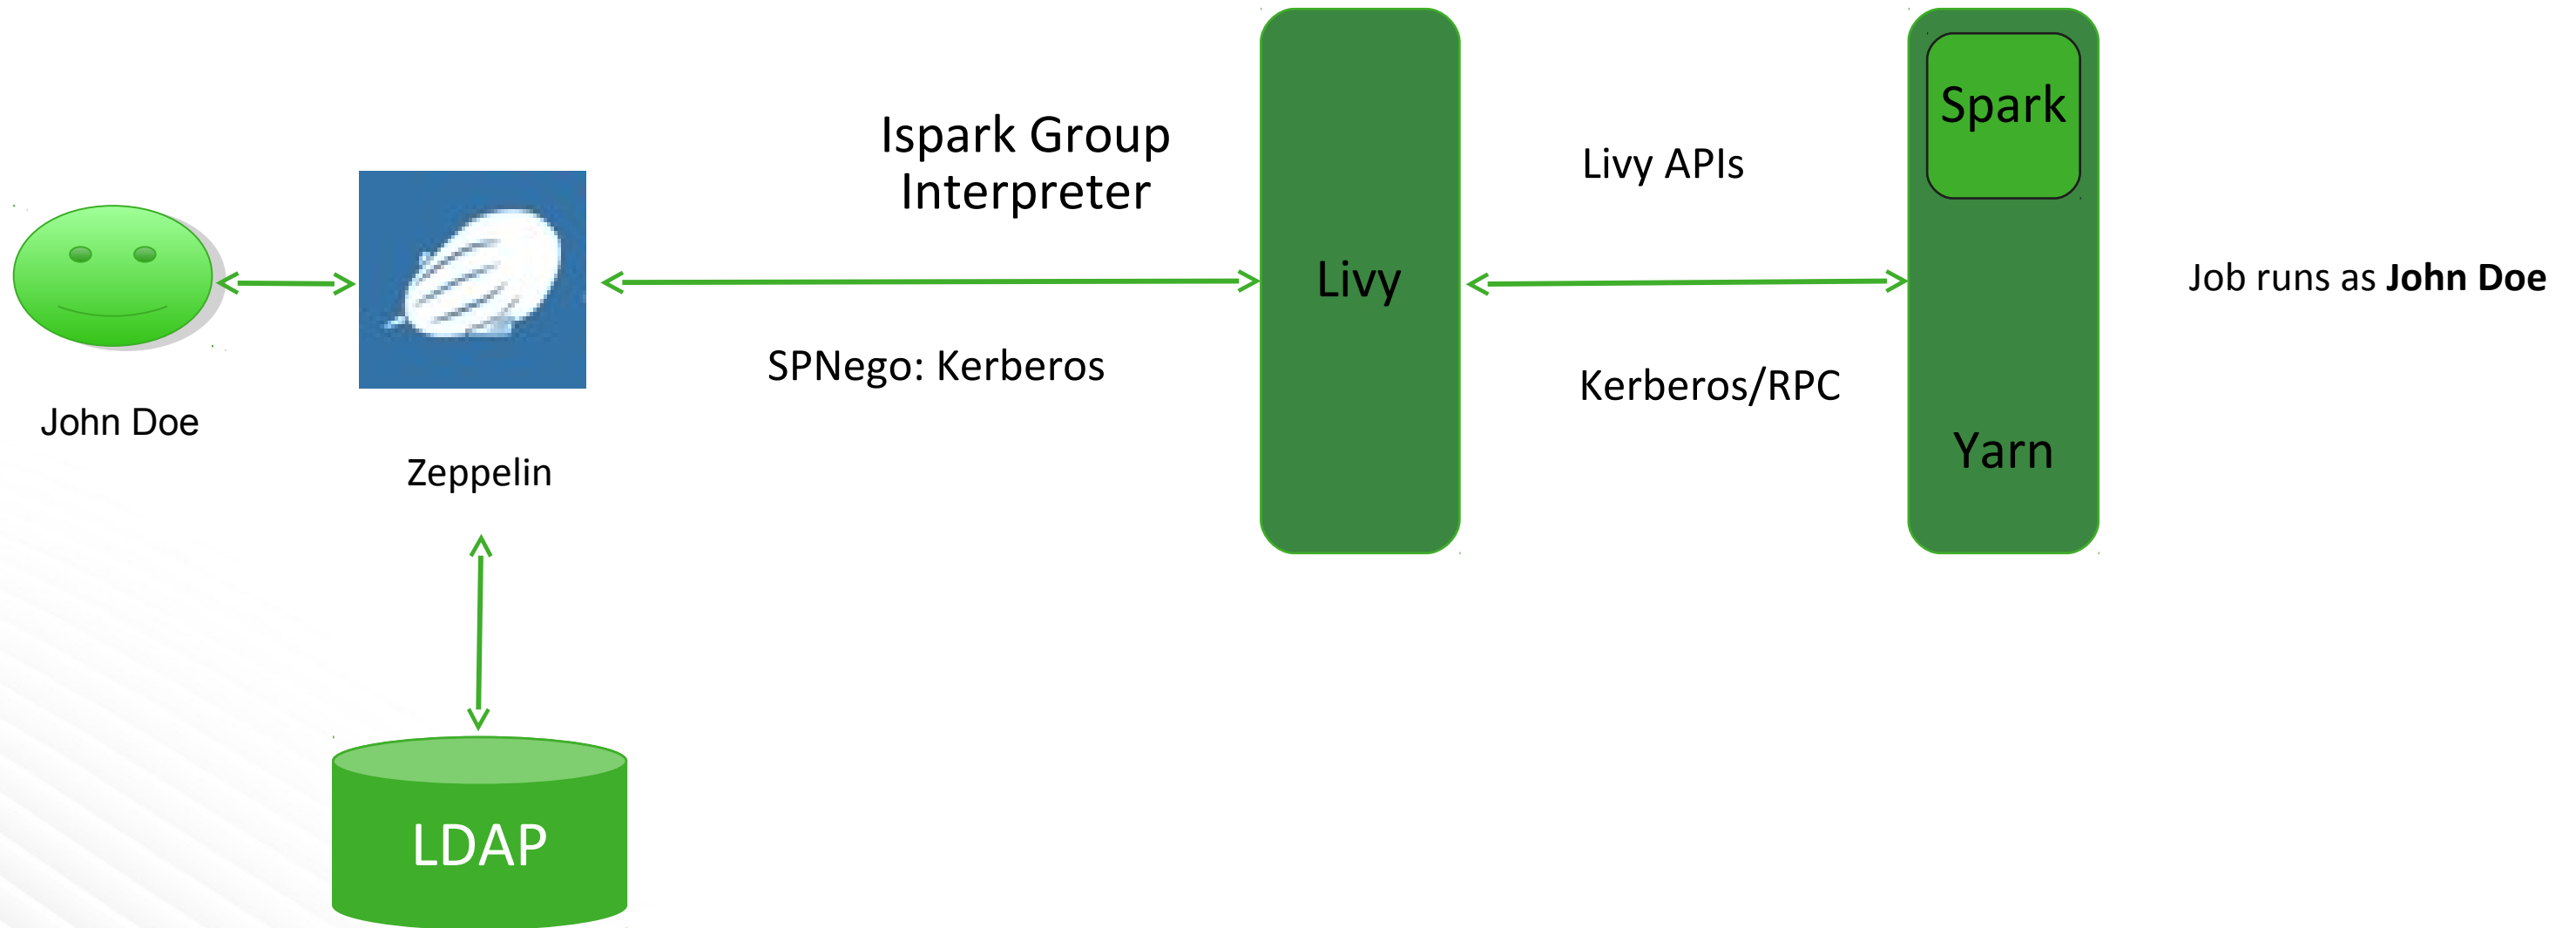


# Apache Zeppelin Security

# Apache Zeppelin: Authentication + SSL



# Zeppelin + Livy E2E Security



# Apache Zeppelin: Authorization

## Authorization in Zeppelin

- Note level authorization
  - Grant Permissions (Owner, Reader, Writer) to users/groups on Notes
  - LDAP Group integration
- Zeppelin UI Authorization
  - Allow only admins to configure interpreter
  - Configured in shiro.ini

```
[urls]
/api/interpreter/** = authc, roles[admin]

/api/configurations/** = authc, roles[admin]

/api/credential/** = authc, roles[admin]
```

## Authorization at Data Level

- For Spark with Zeppelin > Livy > Spark
  - Identity Propagation Jobs run as End-User
- For Hive with Zeppelin > JDBC interpreter
- Shell Interpreter
  - Runs as end-user

# Apache Zeppelin: Credentials

## Credentials in Zeppelin

- LDAP/AD account
  - Zeppelin leverages Hadoop Credential API
- Interpreter Credentials
  - Not solved yet
- Credentials



This is  
still an  
open  
issue



# Apache Zeppelin: AD Authentication

## Configure Zeppelin to Authenticate users

1. /etc/zeppelin/conf/shiro.ini

```
[urls]
```

```
/api/version = anon
```

```
/** = authc
```

Zeppelin leverages  
Apache Shiro for  
authentication/authoriza  
tion

# Apache Zeppelin: AD Authentication

Skip step 1 if securing  
LDAP password is not an  
issue

## Active Directory Authentication

### 1. Create an entry for AD credential

- Zeppelin leverages Hadoop Credential API
- `>hadoop credential create`
- `activeDirectoryRealm.systemPassword -provider jceks://etc/zeppelin/conf/credentials.jceks`
- `chmod 400` with only Zeppelin process r/w access, no other user allowed accessCredentials

### 1. Configure Zeppelin to use AD

```
activeDirectoryRealm = org.apache.zeppelin.server.ActiveDirectoryGroupRealm
activeDirectoryRealm.systemUsername = CN=Administrator,CN=Users,DC=HWQE,DC=HORTONWORKS,DC=COM
#activeDirectoryRealm.systemPassword = Password1!
activeDirectoryRealm.hadoopSecurityCredentialPath = jceks://etc/zeppelin/conf/credentials.jceks
activeDirectoryRealm.searchBase = CN=Users,DC=HWQE,DC=HORTONWORKS,DC=COM
activeDirectoryRealm.url = ldap://ad-nano.qe.hortonworks.com:389
activeDirectoryRealm.groupRolesMap =
"CN=admin,OU=groups,DC=HWQE,DC=HORTONWORKS,DC=COM": "admin", "CN=finance,OU=groups,DC=HWQE,DC=HORTONWORK
S,DC=COM": "finance", "CN=zeppelin,OU=groups,DC=HWQE,DC=HORTONWORKS,DC=COM": "zeppelin"
activeDirectoryRealm.authorizationCachingEnabled = true
```

# Spark Performance

# #1 Big or Small Executor ?

```
spark-submit --master yarn \  
  --deploy-mode client \  
  --num-executors ? \  
  --executor-cores ? \  
  --executor-memory ? \  
  --class MySimpleApp \  
  mySimpleApp.jar \  
  arg1 arg2
```

# Show the Details




Cluster resource: 72 cores + 288GB Memory

--executor-cores	--executor-memory	--num-executor
1	4GB	72
2	8GB	36
3	12GB	24
⋮	⋮	⋮
24	96GB	3



# Benchmark with Different Combinations



# The lower the better

18 cores 108GB memory per node

# Big Or Small Executor ?

- Avoid too small or too big executor.
  - Small executor will decrease CPU and memory efficiency.
  - Big executor will introduce heavy GC overhead.
- Usually 3 ~ 6 cores and 10 ~ 40 GB of memory per executor is a preferable choice.

# Any Other Thing ?

- ❑ Executor memory != Container memory
- ❑ Container memory = executor memory + overhead memory (10% of executory memory)
- ❑ Leave some resources to os and other services

--executor-cores	--executor-memory	--num-executor
3	10GB ( <del>12GB</del> )	24
4	13GB ( <del>16GB</del> )	18
6	20GB ( <del>24GB</del> )	12

- ❑ Enable CPU scheduling if you want to constrain CPU usage<sup>#</sup>

#CPU scheduling -  
<http://hortonworks.com/blog/managing-cpu-resources-in->

# Multi-tenancy for Spark

## Cluster resource Utilization

- Leverage YARN queues
  - Set user quotas
  - Set default yarn-queue in spark-defaults
  - User can override for each job
- Leverage Dynamic Resource Allocation
  - Specify range of executors a job uses
  - This needs shuffle service to be used

# Thank You

**Vinay Shukla**  
 **@neomythos**